

# Git Repositories

## Recommended end-to-end Spark ETL Scala repositories

1. **yennanliu / spark-etl-pipeline** — Demo **ETL/stream & batch examples** (multiple small projects and examples). Good for learning practical ETL flows and many small demos in one repo.

Repository: [GitHub](#)

Quick note: contains multiple subprojects (batch, streaming, EMR demos). To run on Spark 4: update `build.sbt` dependency versions to `4.0.0`, set `scalaVersion := "2.13.16"`, and build with sbt 1.10.x. Watch for plugin/dependency upgrades.

2. **skalskibukowa / Project-Spark-Scala-ETL** — A focused ETL project that extracts from CSV / Postgres, transforms and writes to sinks (**CSV/Parquet/Postgres**). Good for typical ingestion → transform → load pipelines.

Repository: [GitHub](#)

Quick note: likely uses older Spark/Scala versions — update `build.sbt/pom.xml` as above; confirm any JDBC driver compatibility with Java 17.

3. **guruprvn84 / ScalaSparkETLArchetype** — An ETL project template/archetype for Scala + Spark. Useful as a starting skeleton to create production ETL apps with structure & conventions.

Repository: [GitHub](#)

Quick note: project templates often only need dependency updates and minimal code changes to compile on Spark 4 + Scala 2.13.16.

4. **vbounyasit / MyDataFramework** — A reusable Scala ETL framework built on Spark (framework-style rather than a single app). Good if you want to adopt a framework pattern for multiple pipelines.

Repository: [GitHub](#)

Quick note: framework-level code may require more changes when upgrading Spark versions — run tests after bumping Spark artifacts to 4.0.0.

5. **qwshen / spark-etl-framework** — Pipeline-based Spark-SQL transformation framework. Defines readers/transformers/writers as pipeline actors — great for building structured ETL pipelines.

Repository: [GitHub](#)

Quick note: primarily Spark-SQL focused — easier migration to Spark 4 if you keep SQL/DataFrame APIs.

6. **tosun-si / teams-league-airflow-spark-scala-etl** — End-to-end example showing Airflow orchestration, Scala Spark ETL on Dataproc/Serverless, and BigQuery sink. Real-world orchestration + cloud sink example.

Repository: [GitHub](#)

Quick note: this repo is useful if you plan to run ETL in cloud (GCP Dataproc). For Spark 4 check Dataproc runtime compatibility and change job submission parameters accordingly.

7. **bernhard-42 / Spark-ETL-Atlas** — Example project to demonstrate adding lineage (Apache Atlas) to Spark ETL jobs. Good if you need governance/lineage in your ETL.

Repository: [GitHub](#)

Quick note: Atlas integration often uses Spark listeners/APIs; verify listener APIs still match Spark 4 listener hooks.

---

## Spark 4.0.0 / Scala / Java compatibility references

- Spark 4.0.0 release page and notes (Spark 4 is pre-built with Scala 2.13). [Apache Spark+1](#)

---

## Quick migration checklist (how to make a repo Spark-4/Scala-2.13.16/Java-17 ready)

## 1. Update build file

For sbt (`build.sbt`):

```
scalaVersion := "2.13.16"
```

```
libraryDependencies += "org.apache.spark" %% "spark-sql" %  
"4.0.0" % "provided"
```

- 

- For Maven, set Spark 4.0.0 artifacts and `maven-compiler-plugin/source/target` to Java 17.

## 2. sbt & plugins

- Use sbt 1.10.x (`project/build.properties`) and upgrade any sbt plugins (sbt-assembly, etc.) to versions compatible with sbt 1.10.

## 3. Java 17

- Compile with `-source 17 -target 17` / set `javaHome` to JDK 17.
- Ensure any third-party native libs or connectors are Java 17 compatible.

## 4. API changes

- Run `sbt compile / mvn -DskipTests=false test` and fix deprecations or method signature changes. Spark 4 tightened and removed some older APIs — focus on listener, RDD internals and deprecated SQL functions.

## 5. Third-party libs

- Update connectors (JDBC drivers, Hive client, Hadoop client) to versions that support Java 17 and Spark 4.

## 6. Testing

- Add small local smoke tests (run in `local[*]`) and integration tests. Use `spark-submit --master local[*]` to simulate cluster execution.

c

# Which repos to pick depending on your goal

- Want **many small examples** & quick demos → [yennanliu/spark-etl-pipeline](#). [GitHub](#)
- Need a **project skeleton / best-practices template** → [guruprvn84/ScalaSparkETLArchetype](#). [GitHub](#)

- Want **framework-style pipelines** (reader/transformer/writer) → [qwshen/spark-etl-framework](#). [GitHub](#)
  - Building **cloud-orchestrated ETL** (Airflow + Dataproc) → [tosun-si/teams-league-airflow-spark-scala-etl](#). [GitHub](#)
  - Need **lineage/governance** integration → [bernhard-42/Spark-ETL-Atlas](#). [GitHub](#)
- 

1. [stonezhong/spark\\_etl](#)

A generic ETL pipeline framework for Spark applications built in Python. It supports deployment on various cloud Spark platforms like Databricks, AWS EMR, Google Cloud, Azure, and more.

[GitHub Link](#)

2. [jamesbyars/apache-spark-etl-pipeline-example](#)

An example demonstrating robust ETL pipelines with Apache Spark, using Python and PostgreSQL for data retrieval and processing.

[GitHub Link](#)

3. [temcavanagh/Spark\\_ETL\\_Pipeline](#)

A Spark ETL pipeline for predictive modeling feature engineering, implemented with Spark-SQL and PySpark.

[GitHub Link](#)

4. [qwshen/spark-etl-framework](#)

A Scala-based Spark ETL framework using Spark-SQL with pipeline architecture defined via YAML/JSON/XML, suitable for complex workflows.

[GitHub Link](#)

5. [aphp/spark-etl](#)

Contains modules focusing on scalable ETL processes combining Apache Spark, Hive, Solr, and PostgreSQL.

[GitHub Link](#)

6. [SETL-Framework/setl](#)

A simple Spark-powered ETL framework in Scala designed to modularize and structure Spark

ETL projects.

[GitHub Link](#)

7. [aws-samples/amazon-eks-apache-spark-etl-sample](#)

Examples and best practices for running Apache Spark ETL jobs on Amazon EKS with Dockerized Spark applications.

[GitHub Link](#)

These repositories cover a range of **Scala** and **Python** with **Spark ETL projects** that could match your requirements for Spark 4.0.0, Python 3, and Java 17 environments. Let me know if you want detailed exploration or sample code from any specific repo.